

Hot-Swapping MLflow models

A pythonic approach on AWS lambda





Stavros Niafas

ML Engineer | MLflow Ambassador



<https://github.com/sniafas/PyconGR>

Table of contents

- 01 Keep our models fresh w/out downtime
- 02 ML Inference w/ Serverless
- 03 ML lifecycle & MLflow
- 04 AWS λ
- 05 Hot-Swapping!
- 06 Live Demo!

01 Keeping models fresh w/out downtime

Keeping model fresh w/out downtime

Updating Models in production



- Model Drift
- Changes in data / Missing data trends
- Updating a model can cause downtime

Dynamic Updates

- Zero Downtime: Important for high availability services
- Speeding up the ML lifecycle
- Avoids the need for a full re-deployment of application.

Hot-Swapping

- Event driven architecture
- Replace the current model with the updated, while the service is still running
- Pythonic way to achieve that

02 ML Inference w/ Serverless

ML Inference w/ Serverless

Serverless over simplicity 🙌

- Confusing term, servers are there, you just not manage them.
- Zero infra management, self-managed provisioning & auto-scaling
- Ideal for simple inference environments
- Straightforward to implement. Just like a regular Python script

Reduced time-to-market 📈

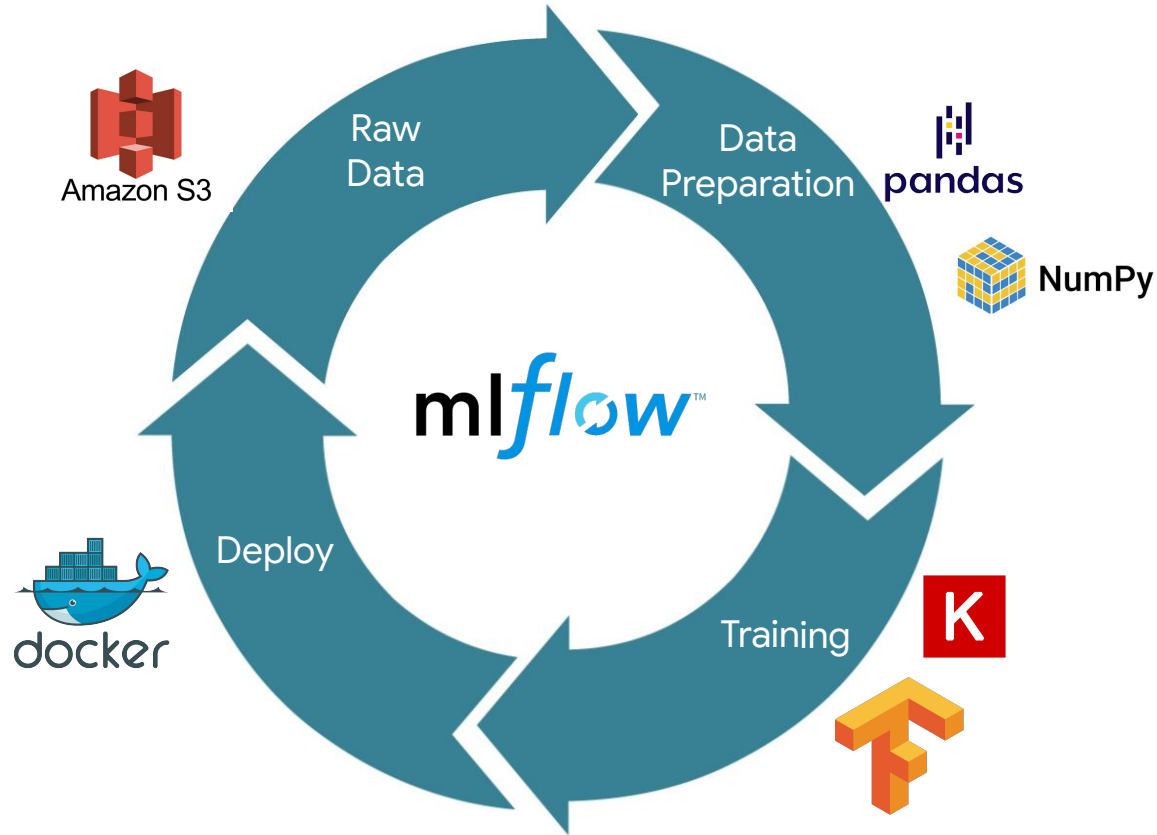
- No overhead for API specifics
- Seamless integration, easy to consume
- Unlimited space for automations and parametrization
- Faster experimentation: Commit, Rebuild, Push

What to look at ⚠️

- Docker packaging & dependencies
- Cold starts 🤖

03 ML Lifecycle & MLflow

ML Lifecycle



MLflow: Managing the ML lifecycle

1. Implement ML lifecycle

MLflow 3.0 Tracking Example [Provide Feedback](#) [Add Description](#)

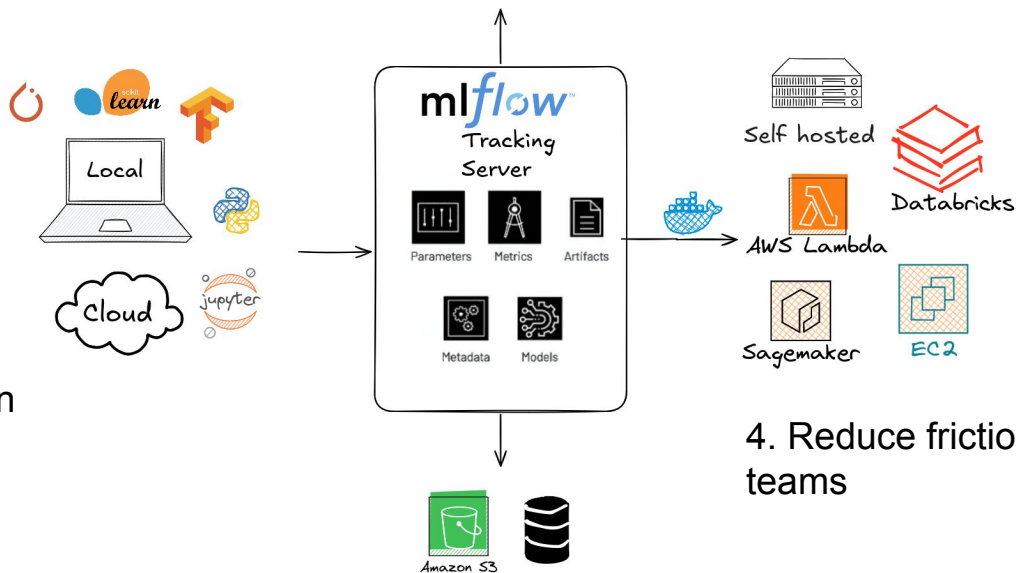
Runs Models Experimental Evaluation Traces

metrics, mse = 0.8

Model attributes			Source run	
Model name	Status	Created	Source run	Model attributes
search-ml-100	Ready	26 seconds ago	popular-unlike-452	
search-ml-40	Ready	29 seconds ago	popular-unlike-452	
search-ml-40	Ready	31 seconds ago	popular-unlike-452	
search-ml-70	Ready	33 seconds ago	popular-unlike-452	
search-ml-40	Ready	35 seconds ago	popular-unlike-452	
search-ml-40	Ready	37 seconds ago	popular-unlike-452	
search-ml-40	Ready	39 seconds ago	popular-unlike-452	
search-ml-50	Ready	41 seconds ago	popular-unlike-452	
search-ml-20	Ready	44 seconds ago	popular-unlike-452	
search-ml-10	Ready	46 seconds ago	popular-unlike-452	
search-ml-0	Ready	48 seconds ago	popular-unlike-452	

3. Governance to track experiments and models efficiently

2. Fluent API in Python



4. Reduce friction between the teams

04 AWS λ

AWS λ

Cold start
initialization



Invocation



```
1 # lambda_function.py
2 import ...
3
4 def get_model():
5
6     model = load_from_s3()
7     return model
8
9 def handler(event, context):
10
11     input_data = json.loads(event["body"])
12
13     model = get_model()
14
15     prediction = model.predict(input_data["data"])
16
17     return {
18         "statusCode": 200,
19         "body": {
20             "prediction": prediction.tolist(),
21         }
22     }
```



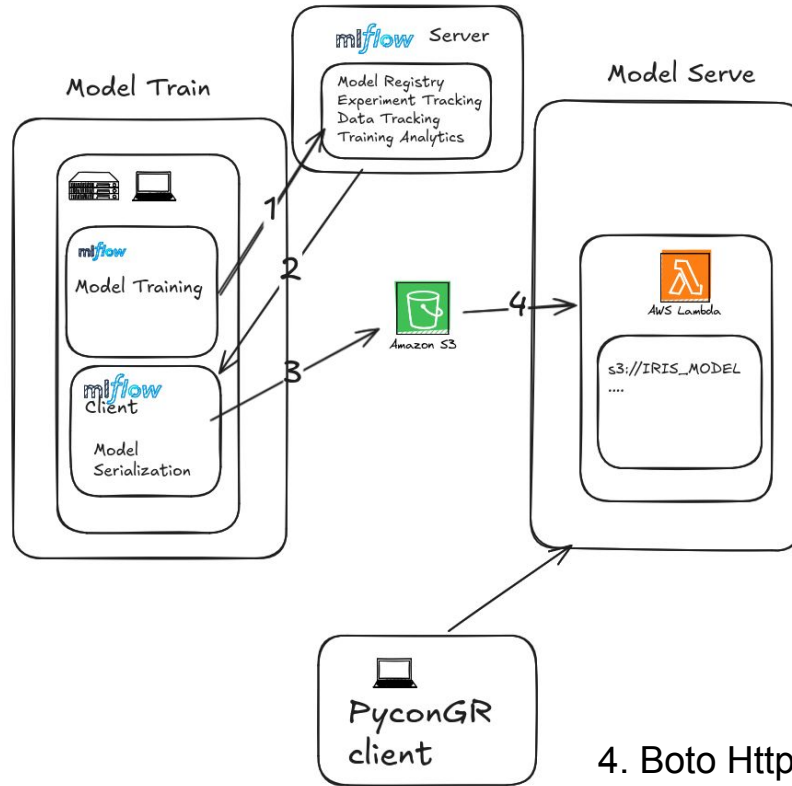
Execution
Environment

05 Hot-Swapping!

Hot-Swapping

1. Model Training
(MLflow)

2. Model
Serialization
(Boto)



3. Model Server
(PyLambda + Boto)

4. Boto Http Client

Hot-Swapping

What's behind zero downtime

Globals



Time check



```
1 model_list, local_timestamp = None, None
2
3 def refresh_model():
4     global model_list, local_timestamp
5
6     # get the current time
7     now = datetime.now()
8
9     # round the time to the minute
10    current_time = now.minute
11
12    # check for updates every N minutes
13    if current_time - local_timestamp > N:
14        new_model_list = load_object_from_s3()
15        local_timestamp = current_time
16        if model_list["version"] != new_model_list["version"]:
17            logger.info("Refreshing models..")
18            load_model_artifacts()
19            model_list = new_model_list
20
21    return model_list
```

05 Live Demo 🙌 🚀



Thank you

